

# Week 1 Assignment

## GPU Intuition & Compute Foundations

GPU Programming using CUDA and Triton (WiDS'25)

Abhineet Agarwal

Dept. of Electrical Engineering, IIT Bombay

## Overview

This assignment is designed to help you understand how GPU architecture differs from CPU architecture, and to analyze a computation-heavy workload from a technical perspective. You will not write CUDA kernels this week – instead, you will build the conceptual and analytical foundation required for effective GPU programming.

## Submission Instructions

Submit a single PDF containing your answers, and include any diagrams as embedded images or separate attachments. Place the completed assignment in the repository folder:

```
week1/  
    assignment1-solution-<your-name>.pdf  
    cpu_baseline-<name>-<task>.py
```

## 1 Task 1: Identify and Analyze a GPU-Accelerable Workload

Choose a computation-heavy workload from your own research, engineering project, or area of interest. Example categories include:

- Simulation loops (finite difference methods, Monte Carlo, N-body, particle updates)
- Numerical kernels (matrix multiplication, reductions, convolutions)
- Machine learning operations (custom layers, loss functions, attention mechanisms)
- Data processing pipelines (pairwise distances, filtering, transforms)

Write a detailed technical analysis (1–2 pages) addressing the following:

### 1.1 Operation Breakdown

- What does the computation do? Describe it precisely.
- Provide pseudocode or a mathematical formulation.
- Specify input sizes, tensor dimensions, or loop bounds.
- Identify independent work units (opportunities for parallelism).

### 1.2 Compute vs Memory Analysis

Using concepts from Week 1 materials:

- Is the operation compute-bound or memory-bound?
- Estimate its approximate arithmetic intensity (qualitative is acceptable).
- Are there reuse opportunities suitable for shared memory?
- Identify dependencies or sequential steps that may limit parallelism.

### 1.3 Expected Behavior on a GPU

- How would CUDA threads map onto the iteration space?
- Would the workload scale well to thousands of threads?
- What challenges might arise (warp divergence, irregular access, small batch sizes)?

### 1.4 CPU Baseline (Can be submitted in Week 1/2/3)

Provide the runtime of a CPU implementation using Python’s `time.perf_counter()` or similar timing tools. This baseline will be used later when implementing the GPU version.

## 2 Task 2: CUDA Execution Model Mapping Diagram

Create a clear diagram illustrating the CUDA execution hierarchy:

Grid → Blocks → Warps → Threads

Annotate your diagram with:

- How your chosen workload's iteration space maps to grid, blocks, and threads
- Where synchronization might be required
- Potential memory bottlenecks
- Where shared memory might be inserted for optimization

You may draw this diagram by hand (scan/photo) or create it digitally.